

Axiom

The automatic data scientist — a five-minute walkthrough to read aloud.

How to use this: Read it aloud, roughly word-for-word — it runs about five minutes at a calm pace. *Italic lines* are delivery cues, not spoken. Slow down, pause where marked, and look up on the bold phrases.

1 · The opening ~30s

Calm, unhurried. Make eye contact before you start.

“Thanks for making the time. In the next five minutes I want to show you something we’ve built called **Axiom** — and why I think it matters for us.

Let me start with a simple question: how long does it take us today to turn a spreadsheet of data into a decision we can trust? Usually the answer is weeks, and it needs a specialist.

Axiom does it in minutes.”

2 · The problem ~45s

“Here’s the problem. We’re sitting on a lot of data — transactions, customers, operations. Inside that data are answers: which transactions are likely fraud, which customers are about to leave, what’s driving our costs.

But getting those answers is hard. It takes a trained data scientist to clean the data, build the model, test it properly, and write it up. That’s expensive, it’s slow, and — the part people underestimate — it’s easy to get subtly wrong. A model can look brilliant in testing and quietly fail in the real world. By then it’s already cost us.”

3 · The solution ~45s

“Axiom solves that. Think of it as an **automatic data scientist** that anyone on the team can use.

You upload a spreadsheet. You point at the one column you want to predict — say, ‘is this transaction fraud?’. You click run. That’s it.

Slow down here — this is the core idea.

Behind the scenes it does everything an expert would: it cleans the data, builds the right signals, trains and compares several models, picks the best one, checks it for mistakes, and hands you a clear report — in plain language, in a few minutes. **No coding, no specialist required.”**

4 · How it works, in one breath ~30s

“Under the hood it’s a pipeline of steps, each doing a specific job and handing off to the next — collect, clean, build features, train, check for errors, improve, and finalise. You don’t need to know any of that to use it. You just see the result: a working model and an honest report.”

5 · Why it matters to the business ~60s

“So why does this matter to us? Three things.

Speed. Work that took weeks now takes minutes. We can test an idea in an afternoon instead of committing a team for a month.

Cost. It runs on our own machines. There’s no per-use AI bill, and it doesn’t need us to hire a specialist for every new question — the expertise is built in.

Trust and privacy. The data never leaves our environment; nothing is sent to an outside AI service. And Axiom is deliberately honest — instead of just claiming a model is ‘accurate’, it tells you, in real terms, how many frauds you’d actually catch for a given amount of review effort. That’s the number a business can plan around.”

6 · Where it shines: fraud ~30s

“We tuned it especially for **fraud detection**, because that’s where the money is and where the mistakes are most costly. Fraud is rare and it changes over time, so it’s easy to build a model that looks good on paper and misses the real thing. Axiom handles those traps properly, and reports results the way a fraud team actually needs to see them.”

7 · The close and the ask ~40s

“The one line I’d leave you with: it’s fast, it’s private, it costs almost nothing to run, and it’s honest about what it can and can’t do. That combination is rare.

Land this next part directly — it’s the ask.

So my ask is simple: let me run Axiom on one of our real datasets — you pick the problem you care about — and I’ll show you a working model and a report by the end of the week. If it delivers, we’ve made data science something the whole team can use, not just a specialist.

Thank you — happy to take any questions.”

If they ask — quick answers

“Does this use ChatGPT / send our data anywhere?”

No. The thinking is classic, proven maths running on our own machines. Nothing leaves our environment, and there’s no per-run AI cost.

“Does it replace our analysts?”

No — it removes the slow, repetitive setup so they spend time on judgement and action, and it lets non-specialists get a first answer on their own.

“How do we know we can trust the model?”

It reports honest, real-world numbers — how many frauds you catch for a given review effort — and it's built to avoid the mistakes that make models look good in testing but fail live.